



CENTER FOR COMPUTATION
& TECHNOLOGY

SAGA

A Simple API for Grid Applications

Introduction to the SAGA API



omii-uk
www.omii.ac.uk



Agenda

- ▣ background recap
- ▣ API structure and scope
- ▣ API walkthrough
- ▣ implementation details
- ▣ API extensions

Grid APIs and Frameworks

- Middleware often targets legacy applications (Unicore, Globus, Condor, ...)
- some are distribution aware (MPICH-G, Ninf-G, . . .)
- few APIs exist for Grid aware applications
 - GridFTP/GRAM
 - DRMAA
 - gLite
 - CoG
 - GAT
 - Cloud APIs

Grid APIs and Frameworks

- ▣ diversity of Grid Middleware implies diversity of APIs
- ▣ some APIs try to generalize Grid programming concepts
- ▣ difficult to keep up with MW development, and to stay **simple**

Open Grid Forum (OGF)

- ▣ The Open Grid Forum (aka GF, EGF, GGF) standardizes distributed computing infrastructures/MW
- ▣ e.g. standardized job description language (JSDL)
- ▣ focuses on interfaces, but also protocols, architecture, APIs

APIs within OGF

- ▣ OGF focuses on services
- ▣ some effort on higher level APIs
 - ▣ Distributed Resource Management Application API (DRMAA)
 - ▣ Remote Procedure Calls (GridRPC)
 - ▣ Checkpoint and Recovery (GridCPR)
 - ▣ Job Submission and Description Language (JSDL)
- ▣ numerous service interfaces, often WS-based (WSRF)

OGF: DRMAA

- implementable on all major resource management services
- simple means to define and submit jobs
- basic job management features (status, kill)
- job templates for bulk job management
- DRMAA.v2 is expected by end of 2010 – OO, extended, SAGA aligned

DRMAA Example

```
drmaa_job_template_t * job_template;

if ( ! ( job_template = create_job_template (exe, 5, 0) ) )
{
    fprintf (stderr, "create_job_template failed\n");
    return 1;
}

while ( ( drmaa_errno = drmaa_run_job (job_id,
                                      sizeof (jobid)-1,
                                      job_template,
                                      diagnosis,
                                      sizeof (diagnosis)-1) )
        == DRMAA_ERRNO_DRM_COMMUNICATION_FAILURE )
{
    fprintf (stderr, "drmaa_run_job failed: %s\n", diagnosis);
    sleep (1);
}
```


OGF: GridRPC

- ▣ standardizes the three existing RPC implementations for Grids
- ▣ example of '*gridified API*'
- ▣ simple: get function handle, call function
- ▣ explicit support for async rpc calls
- ▣ gridrpc.v2 adds support for remote data handles

OGF: GridRPC

```
double A[N*N], B[N*N], C[N*N];

initMatA (N, A);
initMatB (N, B);

grpc_initialize (argv[1]);

grpc_function_handle_t handle;
grpc_function_handle_default (&handle, "mat_mult");

if ( grpc_call (&handle, N, A, B, C) != GRPC_NO_ERROR )
{
    exit (1);
}

grpc_function_handle_destruct (&handle);
grpc_finalize ();
```

OGF: GridCPR

- ▣ Grids seem to favor application level checkpointing
- ▣ GridCPR allows to manage checkpoints
- ▣ defines an architecture, service interfaces, and scope of client API
- ▣ a SAGA extensions exists with a matching client API

OGF: JSDL

- ▣ extensible XML based language for describing job requirements
- ▣ does not cover resource description (on purpose) does not cover workflows, or job dependencies etc (on purpose)
- ▣ JSDL is extensible (ParameterSweep, SPMD, ...)

OGF: JSDL

```
<jsd1:JobDefinition>
  <JobDescription>
    <Application>
      <jsd1-posix:POSIXApplication>
        <Executable>/bin/date</Executable>
      </jsdl-posix:POSIXApplication>
    </Application>
    <Resources ...>
      <OperatingSystem>
        <OperatingSystemType>
          <OperatingSystemName>LINUX</OperatingSystemName>
        </OperatingSystemType>
      </OperatingSystem>
    </Resources>
  </JobDescription>
</jsdl:JobDefinition>
```

OGF: JSDL

- ▣ XML: embeddable into WSRF (WS-Agreement etc.)
- ▣ XML, but relatively flat
- ▣ maps well to existing JDLs, but is 'more complete'
- ▣ extensible (resource description, job dependencies, workflow)
- ▣ top down approach!

OGF: Summary

- some API specs exist in OGF, and are successful
- OGF APIs do not cover the complete OGF scope
- the various API standards are disjoint
- WSDL as service interface specification cannot replace an application level API (wrong level of abstraction)
- **SAGA tries to address these issues**

OGF: top-down vs. bottom-up

- ▣ bottom-up often agrees on (semantic) LCD + backend specific extensions
- ▣ top-down usually focuses on semantics of application requirements
- ▣ bottom-up tends to be more powerful
- ▣ top-down tends to be simpler and more concise
- ▣ ***we very much prefer top-down!***

SAGA

A Simple API for Grid Applications

SAGA

SAGA

Simple API for Grid Applications

SAGA Design Principles

- ▣ SAGA: Simple API for Grid Applications
 - ▣ OGF approach to a uniform API layer (facade)
- ▣ governing principle: 80:20 rule
 - ▣ simplicity versus control!
- ▣ top-down approach
 - ▣ use case driven!
 - ▣ defines application level abstractions
- ▣ extensible
 - ▣ stable look & feel
 - ▣ API packages
- ▣ API Specification is language independent (IDL)
 - ▣ Renderings exist in C++, Python, Java
 - ▣ Examples here are in C++

SAGA Intro: Example 1

```
// SAGA: File Management example

saga::filesystem::directory dir ("any://remote.host.net//data/");

if ( dir.exists ("a") && ! dir.is_dir ("a") )
{
    dir.copy ("a", "b", Overwrite);
}

list <saga::url> names = dir.find ("*-{123}.txt");

saga::filesystem::directory tmp  = dir.open_dir ("tmp/", Create);
saga::filesystem::file      file = dir.open      ("tmp/data.txt");
```

SAGA Intro: Example 1

- ▣ API is clearly POSIX (libc + shell) inspired
- ▣ where is my security??
- ▣ what is 'any:/' ???

SAGA Intro: Example 2

```
// SAGA: Job Submission example

saga::job::description jd;
// details left out

saga::job::service js ("any://remote.host.net/");
saga::job::job      j = js.create_job (jd);

j.run ();

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "RetVal " << j.get_attribute ("ExitCode") << endl;
```

SAGA Intro: Example 2'

```
// SAGA: Job Submission example

saga::job::service js ("any://remote.host.net");
saga::job::job      j = js.run_job ("touch /tmp/touch.me");

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "RetVal " << j.get_attribute ("ExitCode") << endl;
```


SAGA Intro: Example 2

- ▣ stateful objects!
- ▣ yet another job description language? :-(
- ▣ many hidden/default parameters
 - ▣ keeps call signatures small
- ▣ 'any:/' again!
- ▣ TIMTOWTDI (there is more than one way to do it)

SAGA Intro: 10.000 feet


- object oriented: inheritance, interfaces very moderate use of templates though!
- functional and non-functional elements strictly separated
 - non-functional API:
 - look & feel: orthogonal to functional API
 - typically not mappable to remote operations
 - functional API:
 - API 'Packages' extensible
 - typically mappable to remote operations
- few inter-package dependencies - allows for partial implementations

Implementation

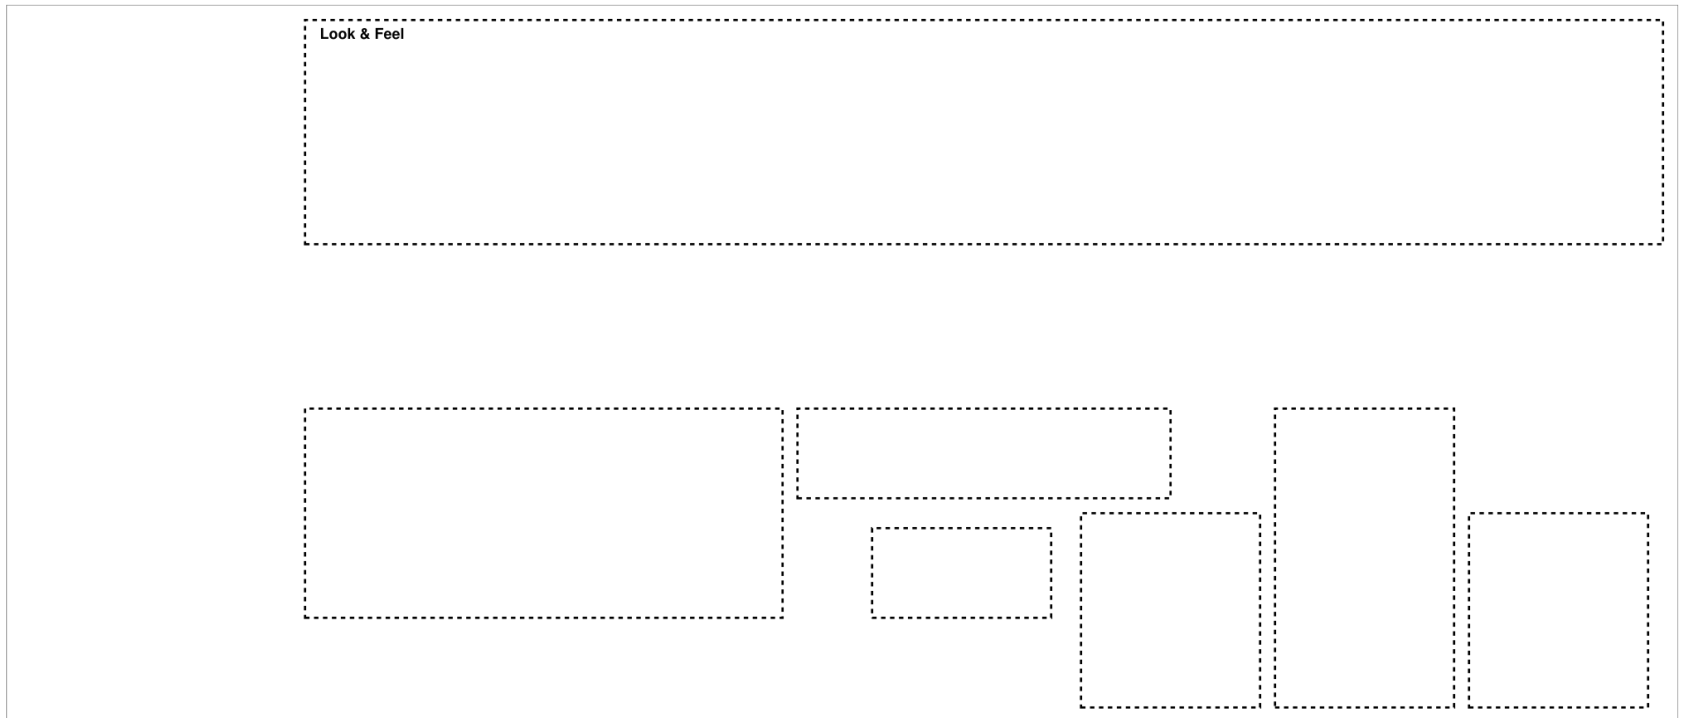


SAGA Class Hierarchy

SAGA: Class hierarchy



SAGA Class hierarchy



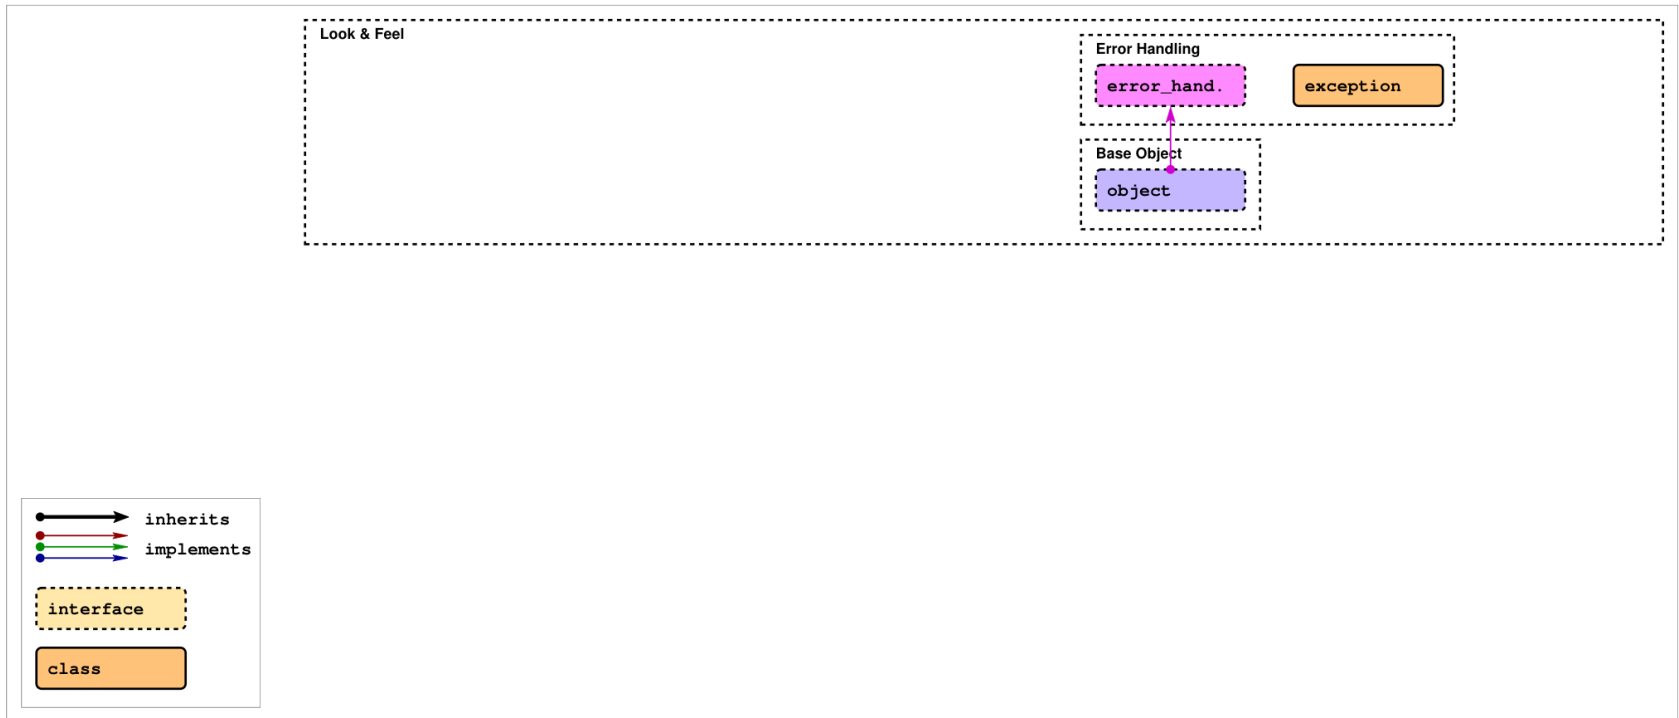
SAGA Class hierarchy



SAGA Look & Feel:

saga::object allows for object uuids, clone() etc.

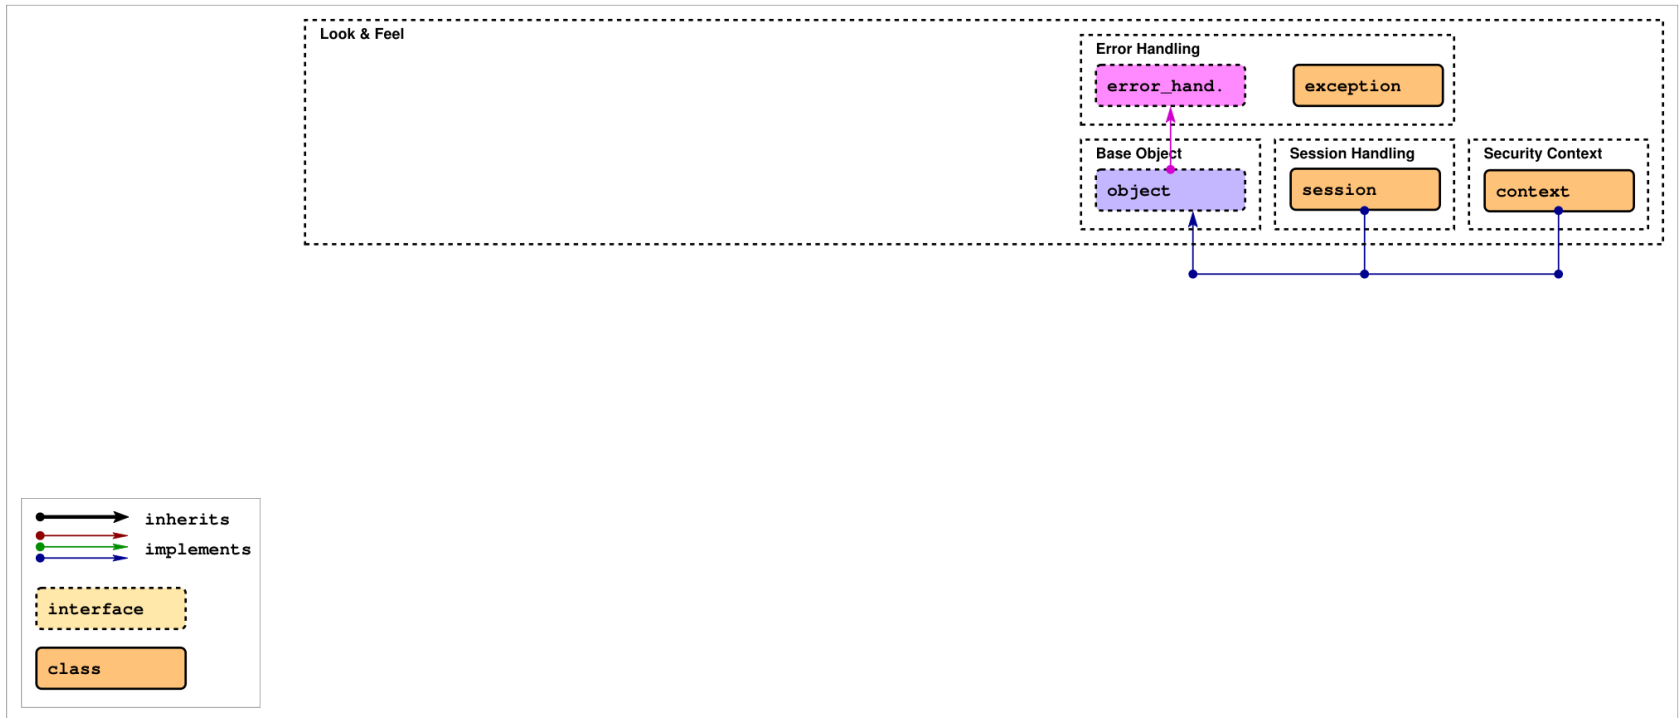
SAGA Class hierarchy



SAGA Look & Feel:

Errors are based on exceptions or error codes.

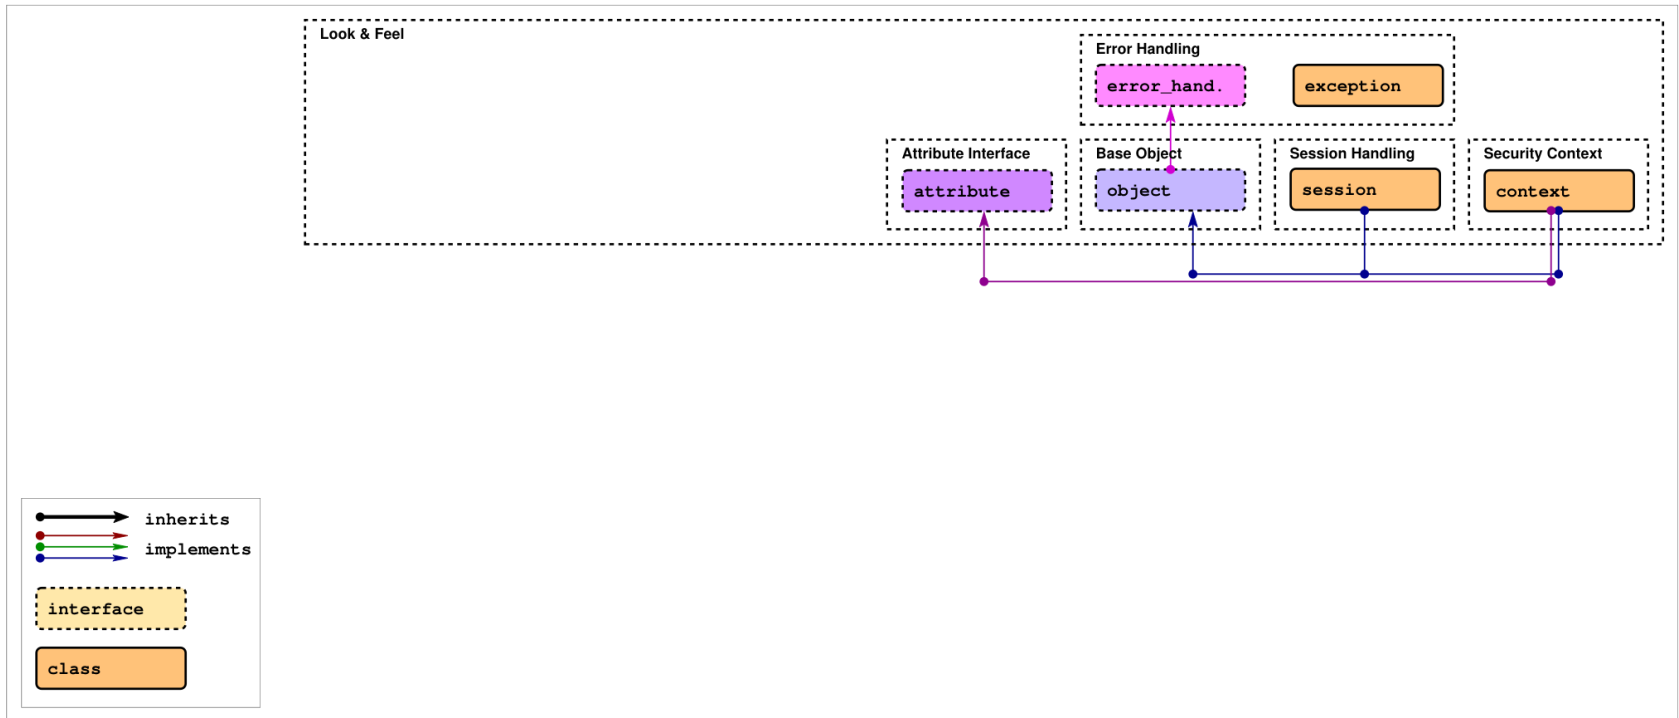
SAGA Class hierarchy



SAGA Look & Feel:

Session and credential management is hidden.

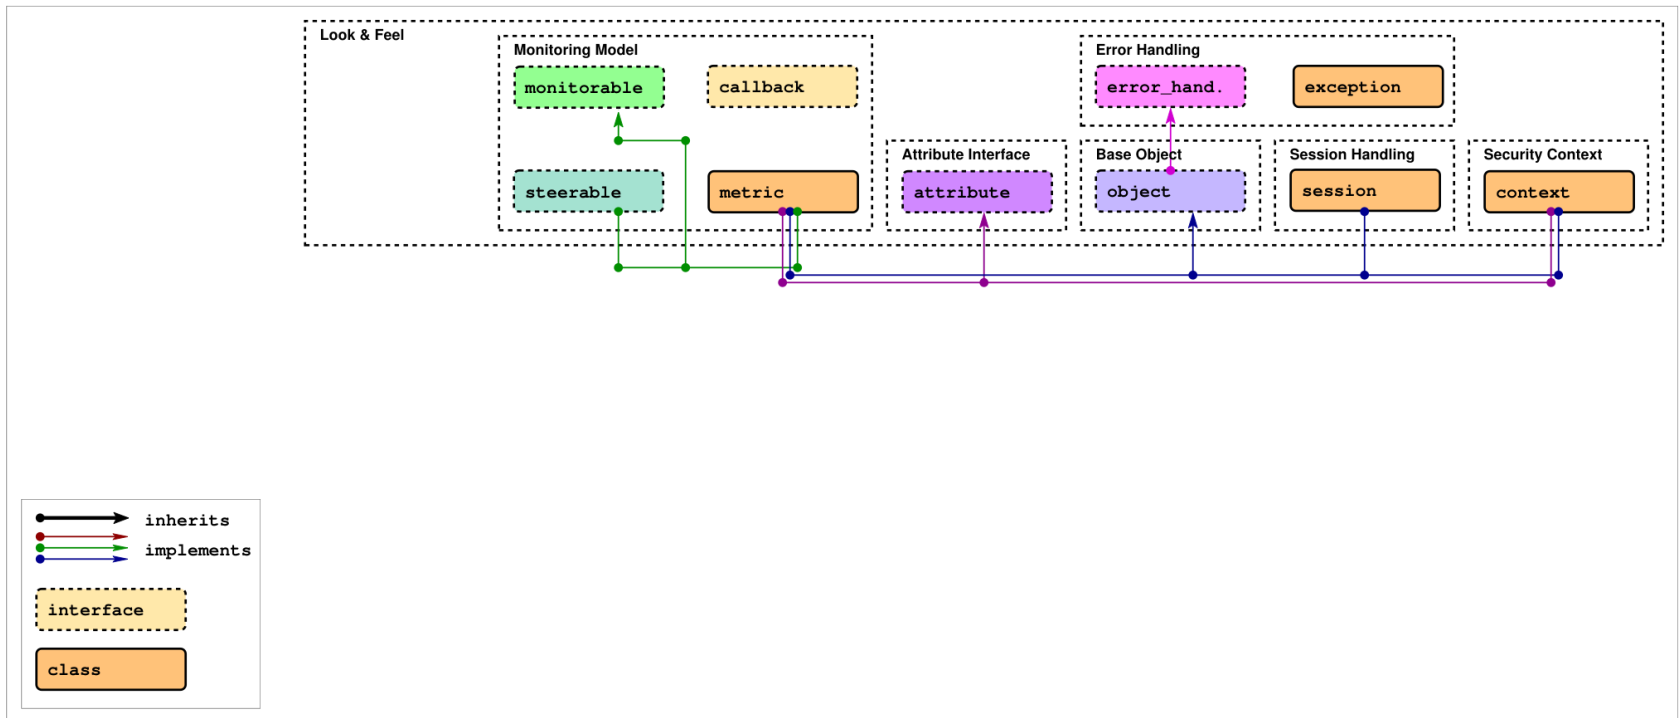
SAGA Class hierarchy



SAGA Look & Feel:

Attribute interface for meta data.

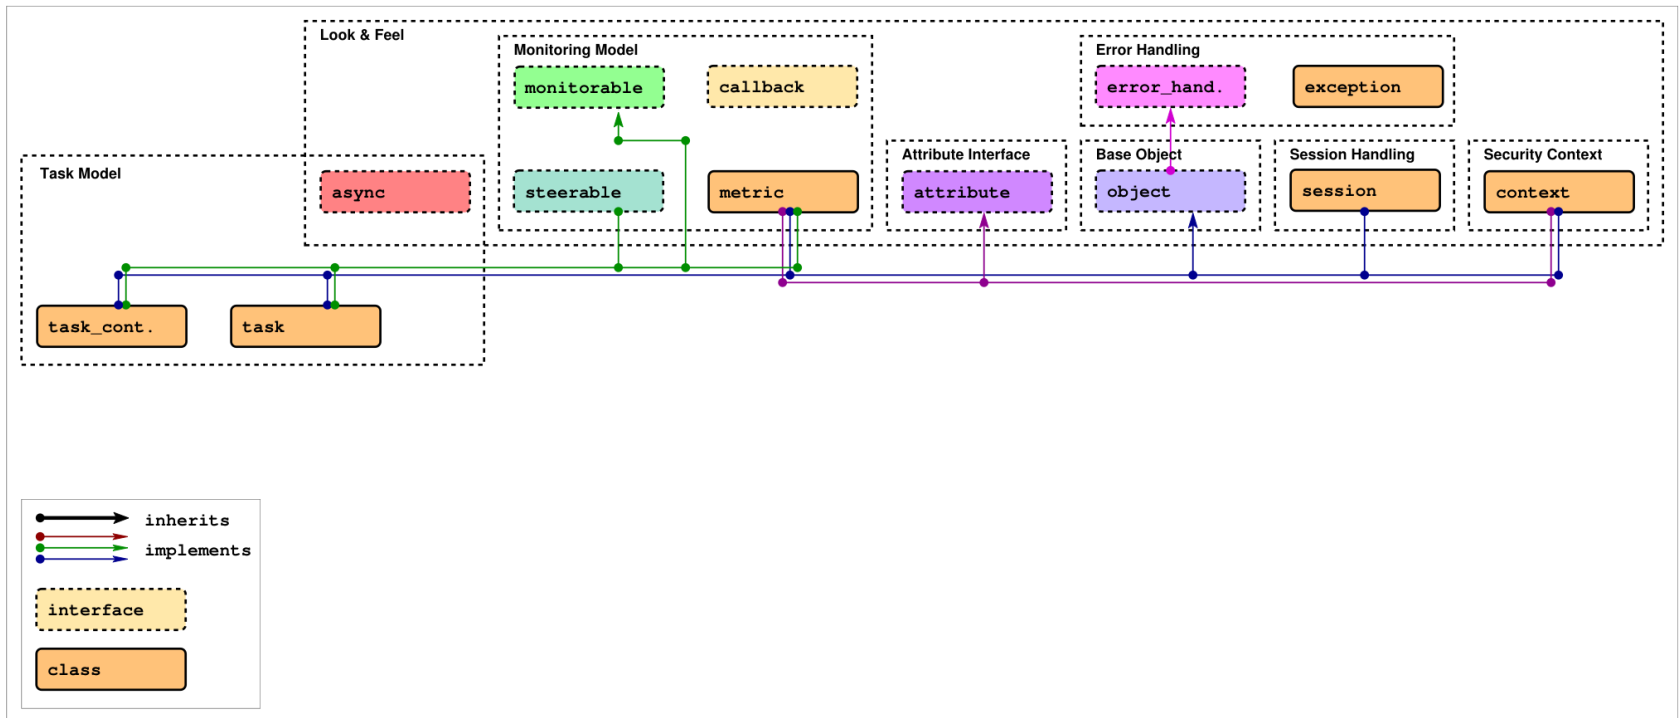
SAGA Class hierarchy



SAGA Look & Feel:

Monitoring includes asynchronous notifications.

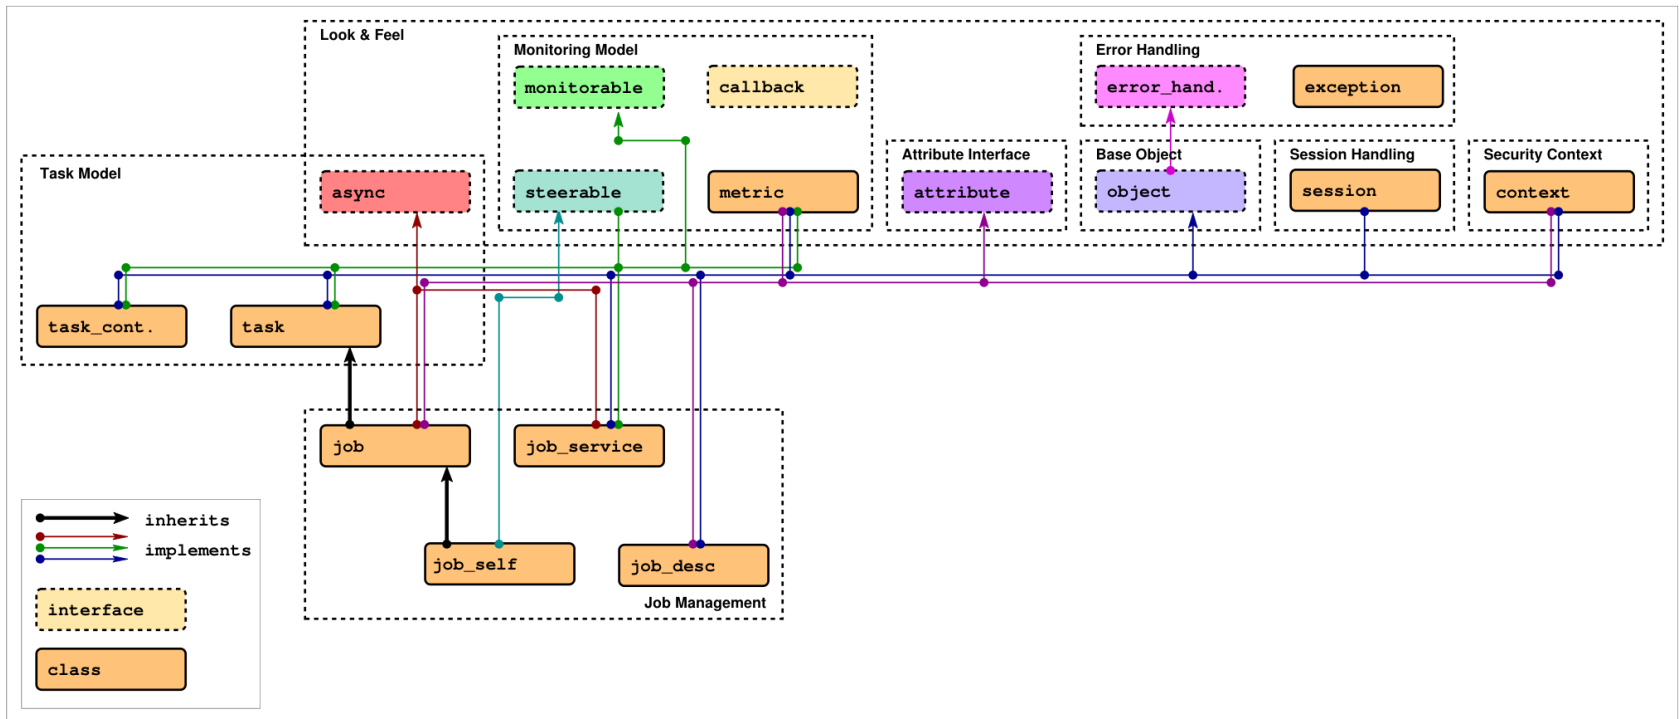
SAGA Class hierarchy



SAGA Look & Feel:

The task model adds asynchronous operations.

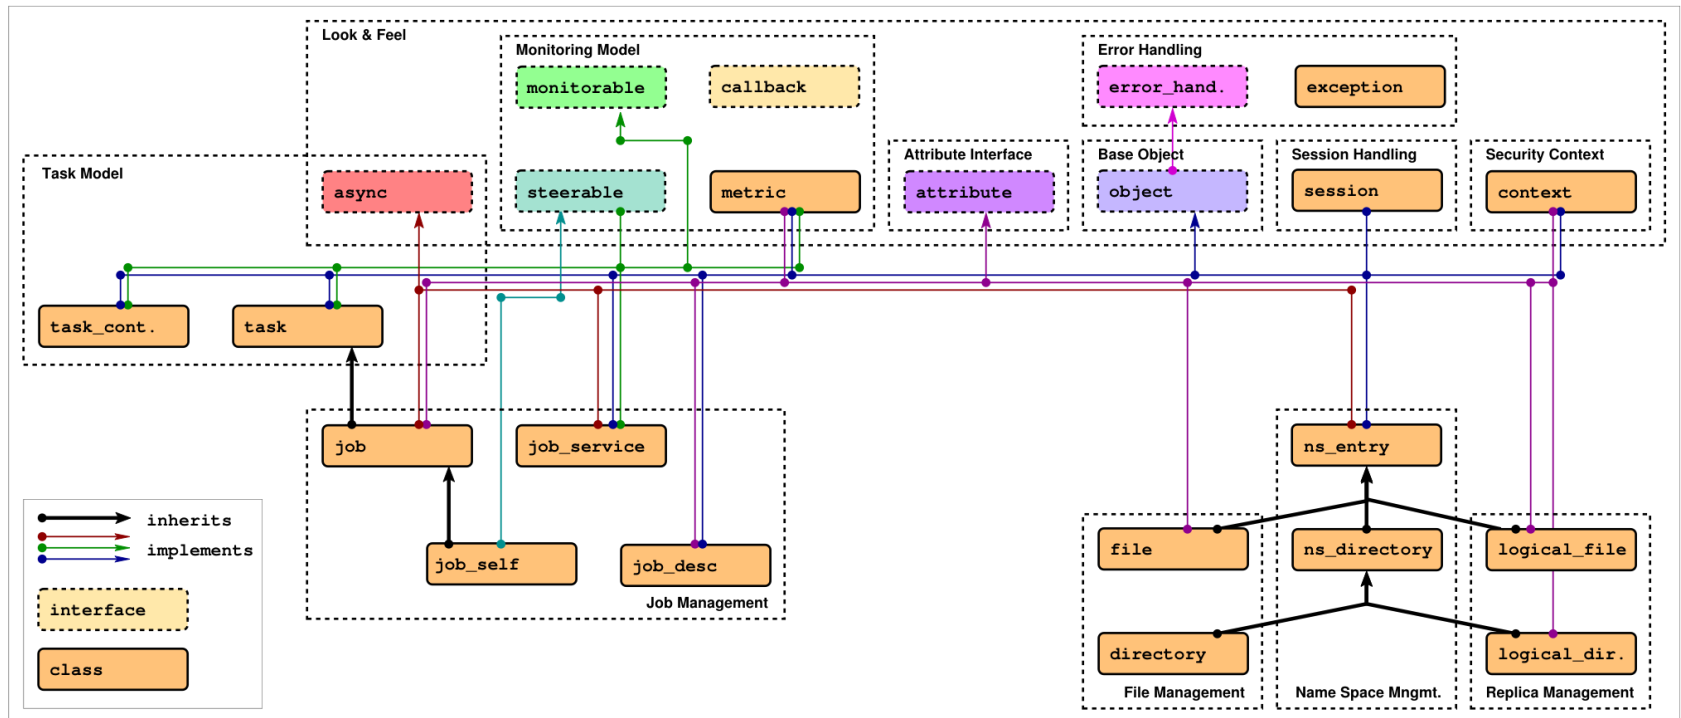
SAGA Class hierarchy



SAGA API Package 'job':

Create and manage remote processes.

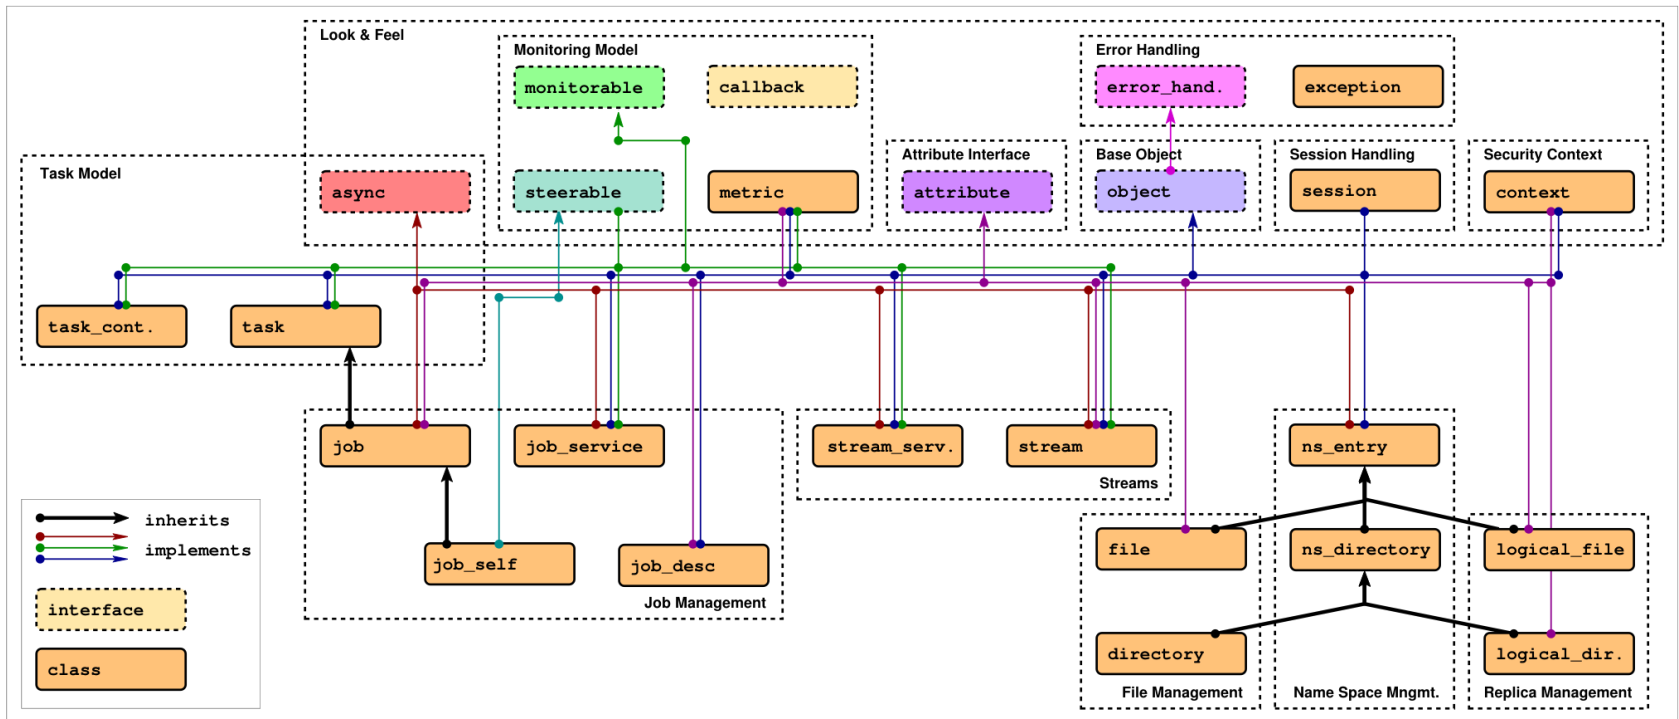
SAGA Class hierarchy



SAGA API Package 'name_spaces':

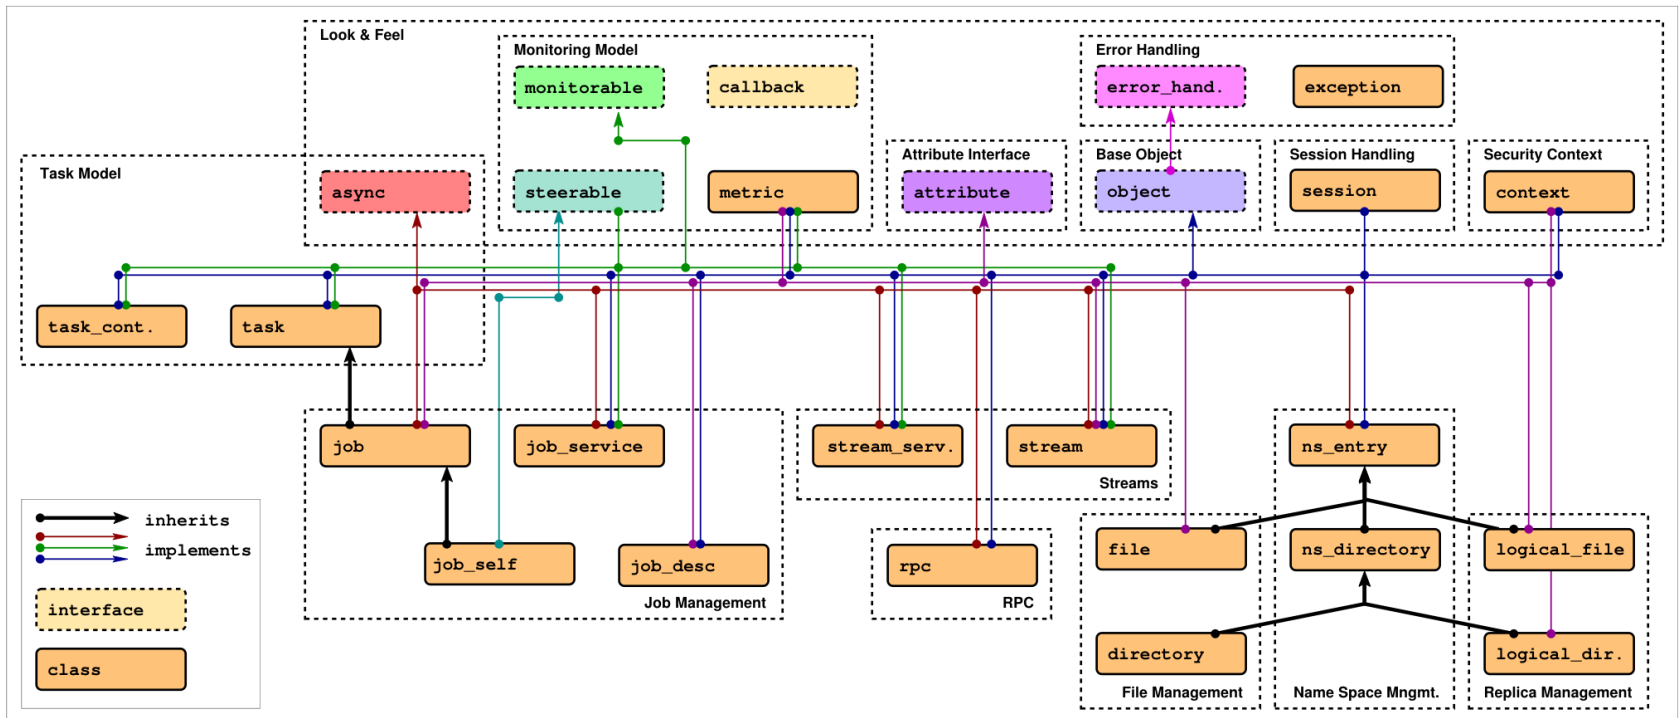
Manage files, replicas, etc.

SAGA Class hierarchy



SAGA API Package 'stream':
SAGA rendering of BSD streams.

SAGA Class hierarchy



SAGA API Package 'rpc':
Remote procedure calls.


SAGA: Class hierarchy

Functional API Packages

SAGA Job Package: Example 1

```
saga::job::description jd;  
saga::job::service      js ("gram://remote.host.net");  
saga::job                j = js.create_job (jd);  
  
j.run ();  
  
cout << "Job State: " << j.get_state () << endl;  
  
j.wait ();  
  
cout << "RetVal " << j.get_attribute ("ExitCode") << endl;
```


SAGA Job Package: job states



SAGA Job Package: job operations

```
j.run      ();  
j.wait     ();  
j.cancel   ();  
  
j.suspend  ();  
j.resume   ();  
  
j.signal    (SIGUSR1);  
j.checkpoint ();  
j.migrate   (jd);
```

SAGA Job Package: job description

```
saga::job::description jd;  
  
jd.set_attribute ("Executable",      "/bin/tail");  
jd.set_attribute ("WorkingDirectory", "data/");  
jd.set_attribute ("Cleanup",         "False");  
  
// pseudo code *blush*  
jd.set_vector_attribute ("Arguments",  ["-f", "my_log"]);  
jd.set_vector_attribute ("Environment", ["TMPDIR=/tmp/"]);  
jd.set_vector_attribute ("FileTransfer", ["my_log >> all_logs"]);
```

SAGA Job Package: job description

Executable	WorkingDirectory	TotalPhysicalMemory
Arguments	<i>Interactive</i>	CPUArchitecture
Environment	Cleanup	OperatingSystemType
CandidateHosts	Input	<i>Queue</i>
SPMDVariation	Output	JobProject
TotalCPUCount	Error	<i>JobContact</i>
NumberOfProcesses	<i>JobStartTime</i>	FileTransfer
ProcessesPerHost	WallTimeLimit	
ThreadsPerProcess	TotalCPUTime	

SAGA Job Package: job description

- ▣ leaning heavily on JSDL, but flat
- ▣ borrowing from DRMAA
- ▣ mixes hardware, software and scheduling attributes!
- ▣ cannot be extended
- ▣ no support for 'native' job descriptions (RSL, JDL, ...)
- ▣ only 'Executable' is required
- ▣ backend MAY ignore unsupported keys
`cd /tmp/data && rm -rf *`

SAGA Job Package: job service

```
saga::job::service js ("gram://remote.host.net/");

vector <string> ids = js.list (); // list known jobs

while ( ids.size () )
{
    string    id = ids.pop_back (); // fetch one job id
    saga::job j  = js.get_job (id); // reconnect to job

    cout << id << " : " << j.get_state () << endl;
}
```

SAGA Job Package: job service

- ▣ represents a specific job submission endpoint
- ▣ job states are maintained on that endpoint (usually)
- ▣ full reconnect may not be possible (I/O streaming)
- ▣ lifetime of state up to backend
- ▣ reconnected jobs may have different job description (lossy translation)

SAGA Job Package: job::self

```
saga::job::service js;  
saga::job::self    self = js.get_self ();  
  
// self represent *this* SAGA application  
  
self.signal        (SIGUSR1);  
self.checkpoint    ();  
self.migrate       (jd);  
  
self.wait          ();    // blocks forever :-P  
self.cancel        ();
```


SAGA Job Package: `job::self`

- ▣ represents the *calling application instance*

- ▣ `self.signal (SIGUSR1);`

goes via the job manager, unlike

`::kill (::getpid (), SIGUSR1);`

- ▣ as it goes over the job manager
(accounting, logging, cleanup...)

- ▣ `job::self` is the only SAGA object to which metrics can be added

SAGA Namespace Package

- ▣ interfaces for managing entities in name spaces
- ▣ files, replicas, information, resources, steering parameter, checkpoints, . . .
- ▣ manages hierarchy (mkdir, cd, ls, . . .)
- ▣ entries are assume opaque (copy, move, delete, ...)

SAGA Namespace Package: example

```
saga::name_space::directory d ("ssh://remote.host.net//data/");

if ( d.is_entry ("a") && ! d.is_dir ("a") )
{
    d.copy ("a", "../b");
    d.link ("../b", "a", Overwrite);
}

list <saga::url> names = d.find ("*-{123}.text.");

saga::name_space::directory tmp = d.open_dir ("tmp/data/1",
                                              saga::name_space::CreateParents);

saga::name_space::entry data = tmp.open ("data.txt");

data.copy ("data.bak", Overwrite); // uses cwd
```

SAGA Namespace Package

- ▣ name space entries are opaque: the name space package can never look inside
- ▣ directories *are* entries (inheritance)
- ▣ inspection: `get_cwd()`, `get_url()`, `get_name()`, `exists()`,
`is_entry()`, `is_dir()`, `is_link()`, `read_link()`
- ▣ manipulation: `create()`, `copy()`, `link()`, `move()`, `remove()`
- ▣ permissions: `permissions_allow()`, `permissions_deny()`
- ▣ wildcards are supported (POSIX influence...)

SAGA Filesystem Package

- ▣ implements name space interface
- ▣ adds access to **content** of namespace::entries (files)
- ▣ POSIX oriented: `read()`, `write()`, `seek()`
- ▣ optimizations: for distributed file access:
 - ▣ scattered I/O
 - ▣ pattern based I/O
 - ▣ extended I/O (from GridFTP)

SAGA Filesystem Package: Example

```
saga::filesystem::file f ("any://remote.host.net/data/data.bin");

char mem[1024];
saga::mutable_buffer buf (mem);

if ( f.get_size () >= 1024 )
{
    buf.set_data (mem + 0, 512);
    f.seek (512, saga::filesystem::Start);
    f.read (buf);
}

if ( f.get_size () >= 512 )
{
    buf.set_data (mem + 512, 512);
    f.seek (0, saga::filesystem::Start);
    f.read (buf);
}
```

SAGA Filesystem Package

- provides access to the **content** of filesystem entries (sequence of bytes)
- saga buffers are used to wrap raw memory buffers
- saga buffers can be allocated/managed by the SAGA implementation
- several incarnations of read/write: posix style, scattered, pattern based

SAGA Filesystem Package: Flags

```
enum flags {  
    None           = 0,  
    Overwrite      = 1,  
    Recursive      = 2,  
    Dereference    = 4,  
    Create         = 8,  
    Exclusive      = 16,  
    Lock           = 32,  
    CreateParents  = 64,  
    Truncate       = 128, // not on name_space  
    Append         = 256, // not on name_space  
    Read           = 512,  
    Write          = 1024,  
    ReadWrite      = 1536, // Read | Write  
    Binary         = 2048 // only on filesystem  
}
```


SAGA Advert Package

- ▣ persistent storage of application level information
- ▣ semantics of information defined by application
- ▣ allows storage of serialized SAGA objects (object persistency)
- ▣ ***very useful for bootstrapping and coordinating distributed application components***

SAGA Advert Package: Example

```
// example for browsing my task adverts
saga::advert::directory todo ("any//remote.host.net/my_tasks/");

// pseudo vector code
list <saga::url> urls = todo.find ("*", ["priority=urgent"]);

while ( urls.size () )
{
    saga::advert ad (urls.pop_front ());
    std::cout << ad.get_attribute ("title")          << std::endl;
    std::cout << ad.get_attribute (« deadline")      << std::endl;
    std::cout << ad.get_attribute ("description") << std::endl << std::endl;
}
```

SAGA Advert Package: Example

```
// master side code: advertise (publish) a saga::file instance
saga::file f (url);

saga::advert ad ("any//remote.host.net/files/my_file_ad", Create);

ad.store_object (f);
```

```
// client side code: retrieve file instance
saga::advert ad ("any//remote.host.net/files/my_file_ad");

saga::file f = ad.retrieve_object ();
```

SAGA: Class hierarchy

Look & Feel Packages

SAGA Session: Example – default session

```
saga::ns_dir dir ("any://remote.host.net//data/");

if ( dir.is_entry ("a") && ! dir.is_dir ("a") )
{
    dir.copy ("a", "../b");
    dir.link ("../b", "a", Overwrite);
}

list <saga::url> names = dir.find ("*-{123}.text.");

saga::name_space::directory sub    = dir.open_dir ("tmp/");
saga::name_space::entry      entry = dir.open      ("data.txt");

entry.copy ("data.bak", Overwrite);
```

SAGA Session: Example – explicit session

```
saga::context c1 (saga::context::X509);  
saga::context c2 (saga::context::X509);  
  
c2.set_attribute ("UserProxy", "/tmp/x509up_u123.special");  
  
saga::session s;  
  
s.add_context (c1);  
s.add_context (c2);  
  
saga::name_space::dir dir (s, "any://remote.host.net/data/");
```

SAGA Session: Properties

- ▣ by default hidden (default session is used)
- ▣ session is identified by lifetime of security credentials, and by objects in this session (jobs etc.)
- ▣ session is used on object creation (optional)
- ▣ `saga::context` can attach security tokens to a session
 - ▣ the default session has default contexts

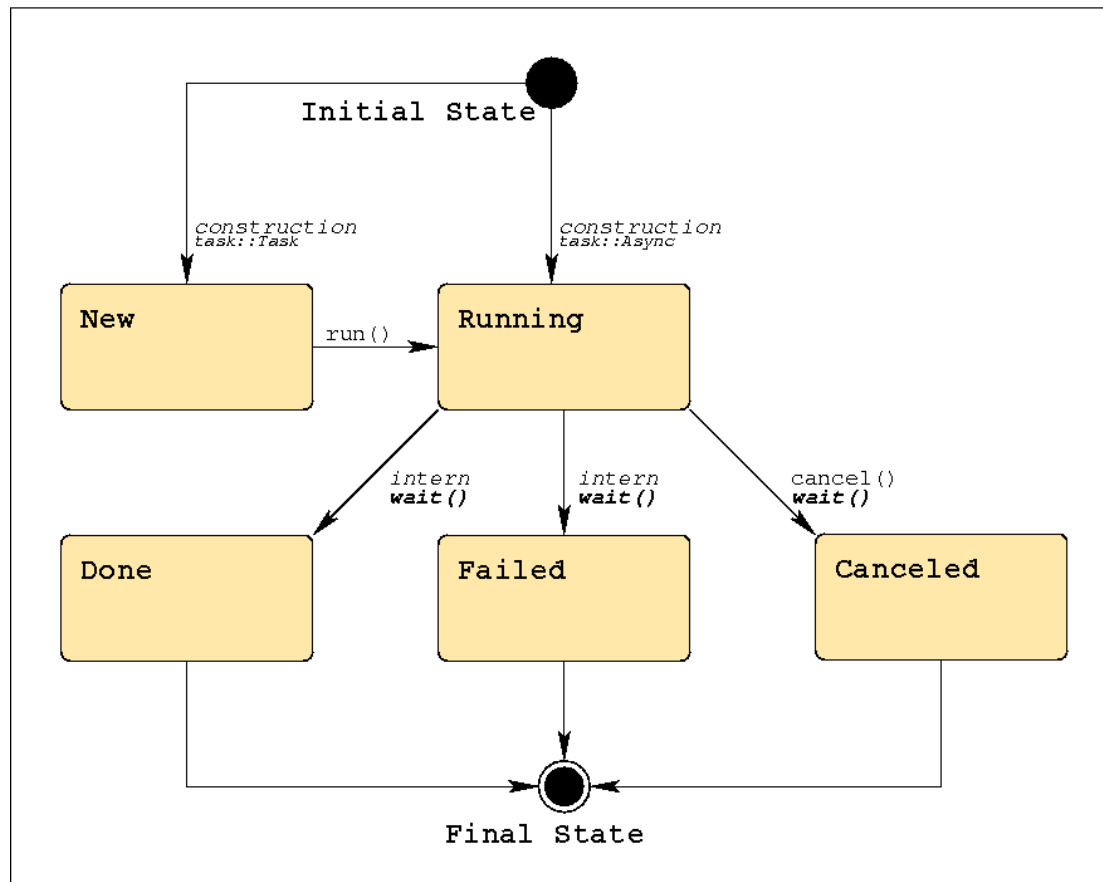
SAGA Session: Lifetime

```
saga::dir  dir (s, "gridftp://remote.host.net/data/");  
  
saga::file file = dir.open ("data.bin");  
  
s.remove_context (c1);  
s.remove_context (c2);  
  
file.copy ("data.bin.bak");  // this works - session is sticky!
```


SAGA Tasks

- ▣ asynchronous operations are a *MUST* in distributed systems, and thus Grids
- ▣ `saga::task` represents an asynchronous operation (e.g. `file.copy ()`)
- ▣ `saga::task_container` manages multiple tasks
- ▣ tasks are stateful (similar to jobs)

SAGA Tasks: States



SAGA Tasks

- ▣ different versions for each method call: Sync, Async, Task
- ▣ signature basically the same
- ▣ differ in state of the task returned by that method
 - ▣ Sync: task is Done
 - ▣ Async: task is Running -> wait();
 - ▣ Task: task is New -> run(); wait();

- ▣ delayed exception delivery

```
if ( saga::task::Failed == task.get_state () )  
{  
    task.rethrow ();  
}
```

SAGA Task: Example

```
// normal method call, synchronous
/* void */ file.copy ("data.bak");

// async versions, never throw (use 'rethrow()' on failure)
saga::task t1 = file.copy      <saga::task::Sync>  ("data.bak.1");
saga::task t2 = file.copy      <saga::task::Async> ("data.bak.2");
saga::task t3 = file.copy      <saga::task::Task>  ("data.bak.3");

// t1: Done
// t2: Running
// t3: New

t3.run ();      // t3 now Running, too

t2.wait ();
t3.wait ();

// t1, t2, t3 are final (Done or Failed)
```

SAGA Task: Example

```
// normal method call, synchronous
size_t s = file.get_size ();

// async versions, never throw (use 'rethrow()' on failure)
saga::task t1 = file.get_size <saga::task::Sync> ();
saga::task t2 = file.get_size <saga::task::Async> ();
saga::task t3 = file.get_size <saga::task::Task> ();

// get_result: implies wait() and rethrow(), and thus can throw!
size_t s1 = t1.get_result <size_t> ();
size_t s2 = t2.get_result <size_t> ();
size_t s3 = t3.get_result <size_t> ();
```

SAGA Task Container: Example

```
// create task container
saga::task_container tc;

// add tasks
tc.add (t1);
tc.add (t2);
tc.add (t3);

// collective operations on all tasks in container
tc.run ();

saga::task done_task = tc.wait (saga::task::Any);

tc.wait (saga::task::All);
```

SAGA Task Container: Tasks and Jobs

```
// NOTE:  
// class saga::job : public saga::task  
// task container can thus manage tasks *and* jobs:  
  
saga::task task = file.copy <saga::task::Async> ("b");  
saga::job job = js.run_job ("remote.host.net", "/bin/date");  
  
saga::task_container tc;  
  
tc.add (task);  
tc.add (job);  
  
tc.wait (saga::task::All);
```